

Statistics and Probability

Joseph Mehdiyev's notes about statistics and probability

Joseph Mehdiyev

If you want to talk about the book(notes) for whatever reason you wish, please contact via my e-mail:
yusifmehdiyev55@gmail.com

Preface

About the Book

This literature is more of like a personal handbook/notebook in the style of the book. The book is for personal use, and is not for commercial purposes. All the source files can be found in the github account, and all the references used are can be found in the reference sections of the each chapter (Also, the general sources are credited below).

The book is just the collection of my notes about statistics and probability. I have organized all of my notes into this piece. This is not supposed to be a proper book.

The book consists of 3 parts. **Probability, Statistical Inference and Statistical Models/Data Mining**. Each section have multiple chapters, references and exercises. Each exercise subsection have **R/Python simulation practices and machine learning examples with visualizations**

To learn the field and write this book, I used various books from known authors, countless mathematics forums about statistics and probability, wikipedia (duh-duh) articles, and MIT OpenCourseWare. These are all the major sources I used while learning the statistics and probability.

- Larry Wasserman - All of Statistics - A Concise Course in Statistical Inference.
- Probability and Statistics by Morris H. DeGroot and Mark J. Schervish.
- R for Data Science.
- An Introduction To Statistical Learning with R/Python (ISLR/ISLP)
- MIT OpenCourseWare Introduction to Probability and Statistics, spring 2022.
- Countless Wikipedia Articles.

The OpenCourseWare has its own license, you can find it here: <https://creativecommons.org/licenses/by-nc-sa/4.0/>

Book's source

This book is fully open source with its source code shared in author's [github](#). You may use the source code for whatever purposes you want to use it for. If you have questions about the book or its source code, you may contact me with my email: yusifmehdiyev55@gmail.com

Code source code

Note that the book is still being written, so you may not find that much code source. In this book, I mainly used **tidyverse** packages to plot, visualize, model, clean data and more. My IDE for R is **RStudio**. You can find R source code in my github, in src directory. For Machine learning and other data science purposes, I used python library **pytorch** and related libraries. I use Neovim/Vim for Python.

Contents

I	Probability	4
1	Introduction to Probability	5
1.1	Set Theory	5
1.2	Probability Law	7
1.3	Probability Distributions	8
1.4	Independent Events	8
1.5	Conditional Probability	8
1.6	Bayes' Theorem	9
1.7	References	10
1.8	Exercises	10
2	Random Variables	13
2.1	Introduction to Random Variables	13
2.2	Distribution Functions c.d.f, p.m.f, p.d.f, p.p.f	14
2.3	Important Random Variables and their distribution	16
2.4	Multivariate Distribution	18
2.5	Marginal Distribution	18
2.6	Independence	19
2.7	Conditioning	19
2.8	Transformations of a Random Variable	19
2.9	References	20
2.10	Exercises	20
3	Expectations and Invariance	22
3.1	Expectation of a Random Variable	22
3.2	Conditional Expectation	24
3.3	Variance	24
3.4	Conditional Variance	25
3.5	Covariance and Corelation	25
3.6	References	26
3.7	Exercises	26
4	Statistical Inequalities	27
4.1	Probability inequalities	27
4.2	Expectation inequalities	28
4.3	References	28
4.4	Exercises	28
5	Convergence, CLT and LLN	29
5.1	Introduction	29
5.2	Types of Convergence	29
5.3	Properties of Convergences	30
5.4	LLN, Law of Large Numbers	31
5.5	CLT, Central Limit Theorem	31
5.6	References	31
5.7	Exercises	32

II	Statistical Inference	33
6	Statistical Models and Statistical Inference	34
6.1	Model and Inference	34
6.2	Point Estimation	34
6.3	Properties of Point Estimation: Efficiency, Consistency, Sufficiency	35
6.4	Confidence Intervals	36
6.5	Hypothesis Testing	37
6.6	References	37
6.7	Exercises	37
7	Methods of Estimation (Parametric Inference)	38
7.1	Method of Moments	38
7.2	Method of Maximum Likelihood	38
7.3	Properties of MLE	39
8	Estimating Statistical Functions (Non-Parametric Inference)	40
8.1	Empirical Distribution Function, e.d.f	40
8.2	Statistical Functionals	40
8.3	References	41
9	Hypothesis Testing and p-value	42
9.1	To-learn and write	42
9.2	Null and Alternate Hypothesis	42
9.3	p-value	43
9.4	Learn and Write missing sections about Hypothesis Testing(*)	43
9.5	References	43
10	Statistical Decision Theory	44
10.1	Risk Function	44
III	Statistical Models	45
11	Linear Regression	46
11.1	Simple Linear Regression Model	46
11.2	Maximum Likelihood and Least Square Estimators	47
11.3	Properties of Least Square Estimators	47
11.4	Multivariate Regression	48
11.5	References	48
12	Classification	50
12.1	Logistic Regression	50
12.2	Estimating Logistic Coefficients	50
12.3	Multinomial Logistic Regression	51
12.4	Estimating Multinomial Logistic coefficients	51
12.5	Linear Discriminant Analysis, l.d.a or LDA	52
12.6	QDA	52
12.7	KNN, write the above chapters after finish learning(*)	52
12.8	References	52
12.9	Exercises	52
13	Resampling Methods	53
13.1	Bootstrap	53
13.2	Basic Introduction	53
13.3	References	53
14	Deep learning	54
15	Unsupervised Learning	55

Part I

Probability

Chapter 1

Introduction to Probability

The concept “probability” is used very often in everyday language to describe the chance of something happening. Mathematically, Probability is a language to quantify uncertainty. This chapter will introduce necessary and basic concepts and namely, **Probability Theory**. We will start the chapter about interpretations of probability.

We will introduce the concepts fast. since they are rather easy compared to next chapters.

Discrete Versus Continuous Concepts

Before we even begin with our concepts. we must learn the difference between the terms **Discrete** and **Continuous** probabilities. Through the book, we will use these terms many times. The Mathematics bluntly can be divided into two distinct categories: *Continuous* mathematics is the study of the objects are uncountable values i.e real numbers, intervals of real numbers and so on; *Discrete* mathematics are study of countable objects. Take probability for example. The probability of simple head and tails experiment is considered discrete, while the probability of weighting 150 kg from intervals 100kg and 200kg is continuous. We will dive deep into these concepts later on, however it is nice to know these terms’ meanings beforehand.

1.1 Set Theory

Set Theory is a branch of mathematics that studies *sets*, which we will define shortly. This branch is, like other parts of mathematics, very deep and complex. We will learn only the most important concepts, which is in high-school level, needed to understand later sections and chapters.

We will quickly introduce the concepts and briefly explain them. The reader may skip this section if they already know about sets and their basic properties.

Sets

A **Set** is a collection of different objects, which are called *elements* of the set. The sets are notated as capital letters such as S . If x is an element of a set S , we write $x \in S$. Otherwise we write $x \notin S$. A set with no elements is called **empty set** and is notated as \emptyset .

If x_1, x_2, \dots, x_n are the elements of the set S , we write:

$$S \in \{x_1, x_2, \dots, x_n\}$$

If S is set of all even numbers smaller than 12, we can draw the diagram as:

We can specify our set as a selection from a larger set. If we want to write the set of all even integers, we can write (Here the set of integers is the universal set):

$$S = \{n \in \mathbb{Z} : \frac{n}{2} \text{ is an integer}\}$$

If a set A ’s elements are also the elements of B , we say that A is a **subset** of B . We can notate it as:

$$A \subseteq B$$

If a set A is subset of B , but is not equal to B , we say that A is **proper subset** of B . We can notate it as:

$$A \subsetneq B$$

Set operations

Union of sets A, B is a set that contains the elements of A and B :

$$A \cup B = \{n : n \in A \vee n \in B\}$$

We can visualize the sets in 2D with circles and their intersections.

Intersection of sets A, B is a set that contains both the elements of A and B :

$$A \cap B = \{n : n \in A \wedge n \in B\}$$

For simplicity we also write $A \cap B = AB$.

Sample Space and Events

The Sample Space, usually denoted as S or Ω , is the *set* of all possible outcomes of an experiment. It is also called **universal set**. Subsets of Ω are called **events**. A sample element of Ω is denoted as ω .

Example 1.1.1. If we toss a six sided dice once, then $\Omega = \{1, 2, 3, 4, 5, 6\}$, the event that the side is even is $A = \{2, 4, 6\}$ while $\omega \in \{1, 2, 3, 4, 5, 6\}$

Example 1.1.2. If we toss a two sided coin twice, then

$$\Omega = \{(HH), (TT), (HT), (TH)\} \wedge \omega \in \{(HH), (TT), (HT), (TH)\}$$

Example 1.1.3. If we toss a 2 sided coin forever, then

$$\Omega = \{\omega = (\omega_1, \omega_2, \dots) : \omega_i \in \{H, T\}\}$$

Example 1.1.4. Let E be the event that only even numbers appear in the six sided dice toss. Then,

$$E = \{2, 4, 6\}$$

With the new definition, we can make more set operation: **complement** of the event A is a set of elements Ω that do not belong to A .

$$A^c = \{n : n \in \Omega \wedge n \notin A\}$$

difference of the set A from B is a set of elements of A that do not also belong to B

$$A \setminus B = A \cap B^c$$

we say that E_1, E_2, \dots, E_N are **disjoint** if

$$A_i \cap A_j = \emptyset$$

A partition of Ω is a sequence of disjoint events such that

$$\bigcup_{i=1}^{\infty} E_i = \Omega$$

Similar to **monotone functions**, we define **monotone increasing** sequence of sets A_1, A_2, \dots as the sequence of sets such that $A_1 \subset A_2 \subset \dots$ and $\lim_{n \rightarrow \infty} A_n = \bigcup A_i$

Moreover, we can define certain rules similar to the rules of algebra:

Commutative laws	$A \cup B = B \cup A$
Associative laws	$(A \cup B) \cup C = A \cup (B \cup C)$
Distributive laws	$A \cap (B \cup C) = (A \cap B) \cup (A \cap C)$

And lastly, **DeMorgan's laws** states that

$$\left(\bigcup_{i=1}^n A_i\right)^c = \bigcap_{i=1}^n A_i^c$$

$$\left(\bigcap_{i=1}^n A_i\right)^c = \bigcup_{i=1}^n A_i^c$$

Which is, in my opinion, very intuitive and can be easily understood with sketching venn diagrams. These are all of the terminology and notations we will be using for learning the probability.

1.2 Probability Law

To show the probability of a event A , we assign a real number $P(A)$ or $\mathbb{P}(A)$ in some textbooks, called **probability of A** . In other words, $P()$ is a unique function with unique properties that inputs an event A , and outputs its probability.

To qualify as probability, P must satisfy 3 axioms:

Axiom 1 $P(A) \geq 0$ for every A

Axiom 2 $P(\Omega) = 1$

Axiom 3 If A_1, A_2, \dots are disjoint:

$$P\left(\bigcup_{i=1}^{\infty} A_i\right) = \sum_{i=1}^{\infty} P(A_i)$$

Let's explain the axioms. The first axiom is very simple, a probability can't be negative, since the meaning of the word probability. Second axiom is also very simple, the probability of any possible outcomes happening is 1, since there must be a outcome at the end of the experiment. Third axiom, assume we have 2 disjoint sets. Then

$$P(A \cup B) = P(A) + P(B)$$

This is true simply because sets are disjoint. Similarly, we can use induction to prove the above property for n sets. Proving for infinite sets are out of scope of this section, therefore we will skip it.

We can derive many properties from these axioms. These are the most simple and intuitive ones:

$$\begin{aligned} P(\emptyset) &= 0 \\ A \subset B &\implies P(A) \leq P(B) \\ 0 &\leq P(A) \leq 1 \\ P(A^c) &= 1 - P(A) \end{aligned}$$

And a less obvious property:

Lemma 1.2.1. For events A and B ,

$$P(A \cup B) = P(A) + P(B) - P(A \cap B)$$

Proof. We can rewrite $A \cup B$ as union of $A \setminus B$, $B \setminus A$, and $A \cap B$, since these are the slices of the thing we want to begin with. Moreover, these slices are disjoint, therefore we can apply our third axiom (P is additive):

$$\begin{aligned} P(A \cup B) &= P((A \setminus B) \cup (B \setminus A) \cup (A \cap B)) \\ &= P(A \setminus B) + P(B \setminus A) + P(A \cap B) \\ &= P(A \setminus B) + P(A \cap B) + P(B \setminus A) + P(A \cap B) - P(A \cap B) \\ &= P(A) + P(B) - P(A \cap B) \end{aligned}$$

□

1.3 Probability Distributions

There are two kinds of Probability Distribution: **Discrete** and **Continuous**. Discrete Probability distribution is the mathematical description of probability of events, that are subsets of **finite or countable infinite** set Ω . If each outcome is equal, then probability of getting 2 even numbers from tossing a six sided dice, which is $\frac{1}{4}$, is an example of this. We can generalize this for event A of finite Ω ,

$$P(A) = \frac{|A|}{|\Omega|}$$

This is the equation almost everybody gets taught in high-school. We can calculate probability of getting heads from tossing a coin, getting a red ball from a box, getting a number from tossing n sided coin and so on. To compute this probability, we first have to count $|\Omega|$ and $|A|$.

For simple experiments, it is rather easy just do count by finger. However, sometimes things get rather complex and we have to use new tools to count them. For example, how many possible outcomes are there from tossing a coin $10^{64^{100}}$ times? We use counting techniques, namely combinatorics. However, the book assumes the reader has knowledge of Combinatorics, therefore we won't introduce the concept here.

Continuous Probability Distribution is similar to its discrete counterpart, however the outcomes are uncountably infinite. Consequently, any probability of selected outcome is 0. Only the events that include these outcomes, making a countable collection of events, have probability themselves.

After we learn about **Random Variables**, we will talk about specific distributions.

1.4 Independent Events

If we flip a six sided dice twice, probability of getting 2 even numbers is $\frac{1}{4}$, which can be found easily just by counting. However, one may guess that we can find the probability for one dice, then square it, which gets the same answer, $\frac{3}{6} \times \frac{3}{6} = \frac{1}{4}$.

This is a prime example of **Independent Events**. The first roll and the second roll are not depended on each other. Whatever the results in first roll can't influence the result in second roll.

The formal definition of independence is,

Definition 1.4.1. Two events A and B are **independent** if

$$P(A \cap B) = P(A)P(B)$$

But how can we know the events are *Independent*? Sometimes, it is rather simple, we know it by logic. Probability of the author being successful is not depended on tossing a coin, it is just simple logic.

In almost all cases, simple logic is enough to determine this property. Another property, is that *disjoint events are never Independent*. Other than that, we have to manually check if the events satisfy the above equation.

Example 1.4.1. Let $A = \{2, 4, 6\}$, $B = \{1, 2, 3, 4\}$. Since $P(A)P(B) = P(AB)$, they are independent.

Example 1.4.2. Let $A = \{2, 4, 6\}$, $B = \{2, 4, 5\}$. Since $P(A)P(B) \neq P(AB)$, they are dependent.

Note that even though $A \cap B \neq \emptyset$, in above examples, the result is not the same. The independence merely shows that another event can't change other event's probability, even though intuitively it makes no sense.

1.5 Conditional Probability

Conditional Probability, as the name implies, is the probability of an event with a condition. More precisely, **Conditional Probability** is the probability of an event A , given that another event B

is already occurred. In such probability, the sample space is reduced to B 's, while we want to find probability of A from B 's space (Which increases of probability of A , since sample space is also reduced). We can show this neatly in venn diagram:

Here are some examples:

Example 1.5.1. If we tossed a six sided dice one time, and we rolled an even number B , what is the probability of getting number 2, event A ?

Since the first toss' result is already happened, we know that $\Omega_{reduced} = \{2, 4, 6\}$ and $A = \{2\}$, then $P(A)_{\Omega_{reduced}} = \frac{1}{3}$.

If there wasn't any condition, the probability of getting 2 would be $\frac{1}{6}$. Simply, in a simple probability we defined a new condition and sort of updated our measurement to $\frac{1}{3}$. This is an important idea in Probability and Statistics, which we will revisit shortly in **Bayes' Rule**

We can show the conditional probability of A given B as:

$$P(A|B) = \frac{P(A \cap B)}{P(B)} \quad \text{for } P(B) \neq 0$$

If we revisit to our simple probability equation, this equation starts making sense since $P(B)$ becomes our reduced sample space, while $P(A \cap B)$ is our event fancily written for condition property.

It is a very common mistake to think $P(A|B) = P(B|A)$, which is easy to understand why just by looking to either venn diagrams or the equations we defined. Moreover, if A and B are independent from each other, then $P(A|B) = P(A)$, which comes from the definition of independence, B can't effect A 's probability.

1.6 Bayes' Theorem

In this section, we will learn about **Bayes' Theorem**, an important concept about probability. This rule is widely used by scientists and programmers. But, what is this rule exactly? Why is it useful?

Bayes' Rule, in simple words, helps to calculate conditional probabilities. It helps us to view probabilities in a degree of belief. I highly recommend watching 3blue1brown's [video](#) about this concept (since visual teaching will always be more practical).

We firstly begin by introducing the simple version of the theorem:

Theorem 1.6.1 (Simplified Bayes' Theorem).

$$P(A|B) = \frac{P(A) \cdot P(B|A)}{P(B)}$$

Proof. We apply the definition of conditional probability twice:

$$P(A|B) = \frac{P(A \cap B)}{P(B)} \quad \wedge \quad P(B|A) = \frac{P(B \cap A)}{P(A)}$$

Using above properties directly gives our theorem. □

Let's try to comprehend the theorem more practically. The theorem can be understood as "Updating the probability of A with a new condition B ". You may think this is an obvious fact and couldn't be that useful. However, let's give some examples that are actually very ambiguous without the theorem.

Example 1.6.1. Steve is a middle aged man living in USA and he is very patient and curious. He also likes debate with people. Which is more likely about Steve: A known mathematician that earned a noble prize or a plumber?

Majority of people would immediately answer "the mathematician", however there is a bigger chance he is a plumber. The reason people get wrong on these questions is because they think that these specific attributes directly corresponds to a smart, wise man. However, they also forget that the number of noble prize winner, middle aged mathematician men that lives in USA is quite low (maybe even zero, I don't really know). The attributes may be

likely to the mathematician, however there is also a low chance that a plumber can have these specific attributes. Also considering there are almost 300k plumbers, the numbers add up.

To not make these kind of mistakes, we must think these attributes, or events as new updates on our main probability, which is a man either being mathematician or a plumber. That is the core idea of Bayes' Theorem.

When using the Bayes' Theorem, it is not always practical to directly calculate the $P(A)$ or $P(B)$. Therefore we need another tool, called **Law of Total Probability** which states that.

Theorem 1.6.2 (Law of Total Probability). Let A_1, A_2, \dots, A_n be partition of Ω . Then for any event B ,

$$P(B) = \sum_{i=1}^n P(B|A_i)P(A_i)$$

Proof. Let $C_i = A_i \cap B$. Then we know that C_1, C_2, \dots, C_n are the partition of B . Therefore using the partition property,

$$P(B) = \sum_{i=1}^n P(C_i) = \sum_{i=1}^n P(A_i \cap B) = \sum_{i=1}^n P(B|A_i)P(A_i)$$

Last step is consequence of conditional probability definition of $P(B|A_i)P(A_i) = P(B \cap A_i)$ \square

This theorem becomes very handy in practical situations. Moreover, with the help of this theorem we can generalize our Bayes' Theorem,

Theorem 1.6.3 (Bayes' Theorem). Let A_1, A_2, \dots, A_n be a partition of Ω such that $P(A_i) > 0$. For $P(B) \neq 0$ and for any $i = 1, 2, \dots, n$,

$$P(A_i|B) = \frac{P(A_i) \cdot P(B|A_i)}{P(B)} = \frac{P(A_i) \cdot P(B|A_i)}{\sum_{i=1}^n P(B|A_i)P(A_i)}$$

Proof. Similar to proof of Theorem 1.7.1, We use definition of conditional probability and lastly apply Theorem 1.7.2 in the last step. \square

1.7 References

1. <https://ocw.mit.edu/courses/18-05-introduction-to-probability-and-statistics-spring-2022/resources/mit18-05.s22.class01-prep-a.pdf/>
2. <https://ocw.mit.edu/courses/18-05-introduction-to-probability-and-statistics-spring-2022/resources/mit18-05.s22.class01-prep-b.pdf/>
3. <http://varianceexplained.org/r/birthday-problem/>

1.8 Exercises

1. 6 card draw.

In a poker game, find the probability of: Two cards have one rank, two cards have another rank, and the remaining two cards have two different ranks.

Solution:

In a 6 card draw, we have total $\binom{52}{6}$ possibilities.

Drawing 2 ranks: $\binom{13}{2}$.

Drawing 2 card from ranks: $\binom{4}{2} \cdot \binom{4}{2}$.

Drawing 2 different ranks again: $\binom{13-2}{2}$.

Drawing 1 card from each ranks: $\binom{4}{1} \cdot \binom{4}{1}$.

Combining these results we get,

$$P(\text{two pair}) = \frac{\binom{13}{2} \cdot \binom{4}{2}^2 \cdot \binom{11}{2} \cdot \binom{4}{1}^2}{\binom{52}{6}} \approx 0.12$$

2. Birthdays: counting and simulation

Consider a group of n people. Consider a set S that is sequence of n birthdays (not necessarily distinct). An example element of S ,

$$\omega = (b_1, b_2, \dots, b_n)$$

- (a) Find the probability of ω .
- (b) Consider 3 events A, B, C
 A : A specific number $\beta \in [1, 365]$ exists in the sequence.
 B : There exists $i, j \in \mathbb{N}$ such that $b_i = b_j$.
 C_k : There exists $i_1, \dots, i_k \in \mathbb{N}$ such that $b_{i_1} = \dots = b_{i_k}$ where $k \leq n$.
 Find probabilities $P(A), P(B), P(C_k)$.
- (c) Calculate probability of $P(C_k)$ for $2 \leq n \leq 50$ via R simulation.

Solutions:

- (a) there are total 365 days. hence,

$$P(\omega) = \frac{1}{365^n}$$

- (b)

$$P(A) = 1 - P(A^c) = 1 - \frac{364^n}{365^n}$$

$$P(B) = 1 - P(B^c) = 1 - \frac{P_{365}^n}{365^n}$$


For C , the exact formula for $P(C_n)$ is complex, since we have to work with cases and such. (That is why simulations are useful).

- (c) This simulation is one of the most standard introductory simulations

```
library(tidyverse)
threshold <- 2:5
ourData <- crossing(trial = 1:3000, people = 1:200) %>%
  mutate(birthday = map(people, ~ sample(365, ., replace = TRUE)),
         multiple = map_int(birthday, ~ max(table(.)))
  )

result <- ourData %>%
  crossing(threshold = 2:5) %>%
  group_by(people, threshold) %>%
  summarize(chance = mean(multiple >= threshold))

result %>%
  mutate(exact = map2_dbl(people, threshold,
                        ~ pbirthday(.x, coincident = .y))) %>%
  ggplot(aes(people, chance, color = factor(threshold))) +
  geom_line() +
  geom_line(aes(y = exact), lty = 2) +
  scale_y_continuous(labels = scales::percent_format()) +
  labs(x = "n",
       y = "Ratio",
       color = "k")
```



Here, the lines with dots show the exact theoretical ratio. As intended, the lines greatly align. We can extract whatever information we need from this exact plot.

Chapter 2

Random Variables

“A random variable can be compared with the holy roman empire: The Holy Roman Empire was not holy, it was not roman, and it was not an empire.” - Some random guy in internet

From the first chapter, we have been using events and sample spaces to develop the idea of probability and calculating it. But, in practical world, we have to link the events and sample spaces to **data**. This concept is called “Random Variable” or r.v shortly. A Random Variable, in informal terms, places the Ω in real line so we can work with it more easily. There are still events, but in terms on Random Variables now.

From now on, we may write Random Variables as r.v or r.v. or r.v.s or r.v or r.v.s shortly for convenience.

2.1 Introduction to Random Variables

A Random Variable describes the data or the outcome ω as a real number. There is a reason this concept exists, since it opens new concepts for practical applications.

Let’s begin with the formal definition of *Random Variable*.

Definition 2.1.1. A **Random Variable X** is a function,

$$X : \Omega \rightarrow \mathbb{R}$$

That assigns a real number $X(\omega)$ to each outcome ω .

In **layman terms**, a random variable is a way to assign a numerical code to each possible outcome. A r.v is **neither random, nor a variable**. It is just a function.

This concept is heavily used instead of sample spaces . From now on, sample space will be mentioned rarely. Think this way, when we work on functions in algebra or sometimes in calculus, we don’t think about about the domain of the function, but the properties of function itself. Here are some examples to understand the concept better.

Example 2.1.1. Flip a coin. We know that $\Omega = \{H, T\}$. A r.v X might assign $X(H) = 1$ and $X(T) = 0$. That is, heads is “coded” as 1 and tails is “coded” as 0.

Example 2.1.2. Flip a fair coin n times. Let X represent the number of heads we get. Then, X is a random variable that takes values $\{0, 1, 2, \dots, n\}$.

Example 2.1.3. Toss a fair six sided dice 2 times. Let X be the sum of the two rolls we get. Then, X is a random variable that takes values $\{2, 3, 4, \dots, 12\}$.

Example 2.1.4. A student wants to write a real number in intervals $[0, 1]$. Let X be the number the student writes. Then, X is also a random variable that takes any real numbers in that interval.

As you may guess, this extremely looks similar to events. Random Variables also have *Independence*, *Conditional Random Variable*, *a probability function* and so on. Additionally, Random Variables can be either **Discrete** or **Continuous**.

Discrete Random Variable's range is finite or countably infinite. The first two examples we gave are Discrete. Continuous Random Variables's range is uncountably infinite like the third example

I want to emphasize that Random Variables are neither random or a variable, they are functions. It is a bit hard to grasp the idea of this concept, so I highly recommend lurking in mathematical forums and try to understand it (that is what I did). But in short, we use Random Variables instead of outcomes, since Random variables are **numbers**. Numbers are easier to work with, we can process the numbers, do algebraic operations to them, also they have a structure that outcomes do not. Turn **Example 2.1.1** to in sample space and events language, which is easier to work with? Bunch of H, T or just a number?

2.2 Distribution Functions c.d.f, p.m.f, p.d.f, p.p.f

c.d.f and p.p.f

We define **Cumulative Distribution Function** as,

Definition 2.2.1. The **Cumulative Distribution Function** or shortly **c.d.f** is a function $F_X : \mathbb{R} \rightarrow [0, 1]$ such that

$$F_X(x) = P(X \leq x)$$

Remark : Every r.v (discrete and continuous) have c.d.f. For this reason, we can use c.d.f for unified treatment of r.v properties (that is, generalized concepts for all r.v). Moreover, c.d.f contains all the information about r.v, both continuous and discrete ones. That is why c.d.f is very useful, even in practical world.

In informal terms, c.d.f is the probability that X will take a value less than or equal to x . This property holds both for continuous and discrete r.v.

Example 2.2.1. We toss a fair coin two times. Let X represent the number of heads we get. Then c.d.f of X is,

$$F_X(x) = \begin{cases} 0 & x < 0 \\ 1/4 & 0 \leq x < 1 \\ 3/4 & 1 \leq x < 2 \\ 1 & x \geq 2 \end{cases}$$

The variable x can get **any real numbers**, such as 2, 4.14 and π . It is just that in discrete case the probability equals to 0 (Which in other words converts discrete to continuous function?). It a bit tricky, they simply take the values from corresponding inequalities. Now, let's look at some properties of c.d.f,

Theorem 2.2.1. Let X have c.d.f F and Y have c.d.f G . If $F(x) = G(x)$ for all x , then,

$$P(X \in A) = P(Y \in A) \quad \text{for all } A$$

Theorem 2.2.2. the function $F : \mathbb{R} \rightarrow [0, 1]$ is a c.d.f for some r.v if and only if F satisfies three conditions:

1. F is non-decreasing

2. F is normalized i.e

$$\lim_{x \rightarrow -\infty} F(x) = 0 \quad \wedge \quad \lim_{x \rightarrow \infty} F(x) = 1$$

3. F is right continuous.

Remarks. The first and second properties are simple and intuitive, therefore we will ignore them (I know it is not the mathematical way, but whatever).

Third property, however, is worth having discussion about. This property directly follows from the inequality \leq . We could even define c.d.f with strict inequality $F = P(X < x)$, and it would still work. It is matter of convention. \square

Definition 2.2.2. Quantile Percent point function, or shortly p.p.f, is defined as inverse of c.d.f i.e,

$$Q(x) = F^{-1}(x)$$

c.d.f and p.d.f

Similar to probabilities of Events, we can calculate probability of X , depending on discrete or Continuous with functions called **Probability Mass Function** and **Probability Density Function**, shortly **p.m.f** and **PDF** respectively,

Definition 2.2.3. If X is discrete, and it takes *countably* values $\{x_1, x_2, \dots, x_n\}$ we define **Probability Mass Function** of X as follows:

$$f_X(x) = P(X = x)$$

Note that $P(X = x)$ is a function, not a number. We have to specify x first to get a number.

Remark: $\{X = x\}$ are disjoint events that form partition of Ω .

With the properties of probability, we have $f_X \geq 0$ for all $x \in \mathbb{R}$ and $\sum_i f_X(x_i) = 1$. Let's revisit our Example 2.2.1

Example 2.2.2. We toss a fair coin two times. Let X represent the number of heads we get. Then c.d.f of X is,

$$f_X(x) = \begin{cases} 1/4 & x = 0 \\ 1/2 & x = 1 \\ 1/4 & x = 2 \\ 0 & \text{otherwise} \end{cases}$$

Moreover, for any set of real numbers, S , we have

$$P(X \in S) = \sum_{x \in S} f_X(x)$$

Since all $\{X = x\}$ are disjoint.

We can apply similar rules to continuous r.vs,

Definition 2.2.4. If X is continuous, we can represent the probability distribution of X with,

$$P(a < X < b) = \int_a^b f_X(x) dx$$

Function f_X is called **Probability Density Function** or PDF as shortly.

Nothing new here really, we just change the properties of p.m.f that we can use it on continuous r.vs. Now, let's look at some examples,

You may noticed that c.d.f is similar to p.m.f and PDF. Indeed, they are related, c.d.f is just sum of these functions we defined over some interval x .

Definition 2.2.5. c.d.f is related to p.m.f and PDF. For discrete r.vs,

$$F_X(x) = P(X \leq x) = \sum_{x_i \leq x} f_X(x_i)$$

And for continuous r.vs,

$$F_X(x) = \int_{-\infty}^x f_X(x) dx$$

And $f_X(x) = F'_X(x)$ for for all differentiable points x .

Note that this definition is heavily used instead of direct definitions above, since we can work with c.d.f only, and derive it to get needed functions.

2.3 Important Random Variables and their distribution

Definition 2.3.1. If X has distribution A , we write

$$X \sim A$$

Usually A depends on some fixed numbers to define properly, we call them **parameters**. For example, the distribution **Bernoulli**(p) has parameter p . We show parameters in c.d.f and p.m.f as,

$$f(x; \text{parameters}) \quad \text{and} \quad F(x; \text{parameters})$$

There are some specific examples of r.v. that are very useful in practical applications. We will show most important ones, and briefly discuss them. In later chapters, we will learn more about them. Note that we will write the notation with the name of the distribution.

Degenerate distribution or Point mass distribution: $X \sim \delta_a$

Consider tossing coin or dice where all the sides show the same value. The p.m.f is ,

$$f_X(x; \delta_a) = 1 \quad \text{for } x = a$$

You might guess why it is called degenerate sometimes. It is not random, but the distribution satisfies the definitions!

Discrete Uniform distribution

This distribution is the one of the most known ones. When there are finitely many values and each of them have the same probability, then $p = \frac{1}{n}$. Simple coin tossing, dice rolling are prime example of these. The p.m.f is,

$$f_X(x) = \frac{1}{n}$$

Where $x \in \{1, 2, \dots, n\}$. Nothing new here. for other cases, $f_X(x) = 0$.

Bernoulli distribution: $X \sim \text{Bernoulli}(p)$

This distribution describes “Yes or No” type of experiments such as coin flipping. Therefore, $P(X = 1) = p$, $P(X = 0) = 1 - p$. We can also calculate p.m.f,

$$f_X(x; p) = p^x (1 - p)^{1-x} \quad \text{for } x \in \{0, 1\}$$

Binomial distribution: $X \sim \text{Binomial}(n, p)$

This distribution is generalized form of **Bernoulli distribution**. Similar to Bernoulli, this distribution describes “Yes or no” type of experiments, but for n times of tries e.g tossing a coin n times. Assuming tries are independent of each other, we can show p.m.f as,

$$f_X(x; n, p) = \binom{n}{x} p^x (1 - p)^{n-x} \quad \text{for } x \in \{0, 1, \dots, n\}$$

Notice that $\text{Binomial}(1, p) = \text{Bernoulli}(p)$.

Geometric distribution: $X \sim \text{Geom}(p)$

This distribution is also specified with Bernoulli. The geometric distribution describes the probability of the first occurrence of success requires after x independent trials e.g getting the first head after x tosses. The p.m.f is,

$$f_X(x; p) = (1 - p)^{x-1} p \quad \text{for } x \geq 1$$

Poisson distribution: $X \sim \text{Poisson}(\lambda)$

This distribution is mainly used for counts of events like photons hitting a detector in a time interval, number of car accidents, students achieving a low and high mark on exam, or number of pieces of chewing gum on a tile of a sidewalk. its p.m.f is,

$$f_X(x; \lambda) = e^{-\lambda} \frac{\lambda^x}{x!} \quad \text{for } x \geq 0$$

Usually, $\lambda = rt$ where r is average rate the events occur and t is the time interval. The r.v X represents the number of events.

Uniform distribution: $X \sim \text{Uniform}(a, b)$

The p.d.f of X is defined as,

$$f_X(x; a, b) = \frac{1}{b - a} \quad \text{for } x \in [a, b]$$

Normal (Gaussian) distribution: $X \sim N(\mu, \sigma^2)$ or $X \sim \mathcal{N}(\mu, \sigma^2)$

This distribution is one of the most popular ones even for non-mathematician, layman people. The famous IQ graph, badly made “memes” are example of this. This distribution plays important role in statistics and probability. Moreover, we can observe this distribution in nature. p.d.f is defined as,

$$f_X(x; \mu, \sigma^2) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}} \quad \text{for } x \in \mathbb{R}$$

We will learn about μ and σ in later chapters. For now, just assume they are some random parameters.

We define **standart Normal distribution** as $N(0, 1)$, r.v as Z . This specific distribution is very important, so much that we show its p.d.f and c.d.f with new notation, namely $\phi(z)$ and $\Phi(z)$. There is no closed form expression for $\Phi(z)$. In modern days, the programming libraries calculate them.

It can be shown that we can show **any normal probabilities** we want with $\Phi(z)$.

Exponential distribution: $X \sim \text{Exp}(\lambda)$

This distribution is continous analogue of the geometric distribution. This distribution (and the ones after this) has complex properties so we will be brief with this. We define its c.d.f as ,

$$f_X(x; \lambda) = \lambda e^{-x\lambda}$$

Gamma distribution: $X \sim \text{Gamma}(\alpha, \beta)$

First, we start with a definition,

Definition 2.3.2. For $\alpha > 0$, we define **Gamma function** as,

$$\Gamma(\alpha) = \int_0^{\infty} t^{\alpha-1} e^{-t} dt$$

It is generalized form of more simple and specific version,

$$\Gamma(n) = (n - 1)! \quad \text{for } n \in \mathbb{N}$$

We define p.d.f of x as,

$$f_X(x; \alpha, \beta) = \frac{1}{\beta^\alpha \Gamma(\alpha)} x^{\alpha-1} e^{-x/\beta} \quad x > 0$$

Notice that $\text{Gamma}(1, \beta) = \text{Exp}(\beta)$. That is, exponential distribution is a specific case of gamma distribution. Gamma distribution itself is very useful and heavily used in this field. Moreover, we can derive more advanced (which I have difficulties understanding) distributions from the gamma function. For now, we will end our discussion here.

2.4 Multivariate Distribution

In practical word, we often work with multiple r.v in the same experiment. This can be a medical research with multiple tests, where tests are related with each other with the same sample space Ω and the same probability.

First, we define a special vector,

Definition 2.4.1. Let X_1, X_2, \dots, X_n be r.vs. We call $X = \{X_1, X_2, \dots, X_n\}$ a **random vector**.

Definition 2.4.2. If r.vs X_1, X_2, \dots, X_n are **independent** and have the same **marginal distribution** with c.d.f F , we define these r.vs as **independent and identically distributed**, shortly i. i. d, with notation,

$$X_1, \dots, X_n \sim F$$

similarly, we show the p.d.f the same way. i.i.d property is very important in statistical field.

We can apply multivariate c.d.f as

Definition 2.4.3. For n r.v $\{X_1, X_2, \dots, X_n\}$, the multivariate c.d.f F_{X_1, X_2, \dots, X_n} is given by,

$$F_{X_1, X_2, \dots, X_n}(x_1, x_2, \dots, x_n) = P(X_1 \leq x_1, \dots, X_n \leq x_n)$$

There is nothing fancy here, actually. We simply redefine c.d.f in general sense for n r.vs.

Similarly, we can define multivariate p.m.f as,

Definition 2.4.4. For random vector X , the multivariate p.m.f f_{X_1, X_2, \dots, X_n} is given by,

$$f_{X_1, X_2, \dots, X_n}(x_1, x_2, \dots, x_n) = P(X_1 = x_1, \dots, X_n = x_n)$$

This is generalized form of p.m.f

Similarly, we define,

Definition 2.4.5. We know that c.d.f and p.d.f are related by derivative. Then, For random vector X , the multivariate p.d.f f_{X_1, \dots, X_n} is given by,

$$f_{X_1, X_2, \dots, X_n}(x_1, x_2, \dots, x_n) = \frac{\partial^n F_{X_1, \dots, X_n}(x_1, \dots, x_n)}{\partial x_1 \partial x_2 \dots \partial x_n}$$

The Properties and theorems are similar, but are generalized for n r.vs.

2.5 Marginal Distribution

If more than one variable is defined in an experiment, it is important to distinguish between the multivariate probability of (X_1, X_2, \dots, X_n) and individual probability distributions of X_1, X_2, \dots, X_n

Formally, **Marginal distribution is the probability of a single event (or r.v) occurring, independent of other events**. Therefore implementing marginal distributions are rather easy. In multivariate distributions, we redefine the needed variable as a “constant” and work with other variables only.

Definition 2.5.1. If X is a random vector with p.m.f f_{X_1, X_2, \dots, X_n} , then we define marginal distribution as,

$$f_{X_1} = P(X_1 = x_1) = \sum_{x_1 \text{ constant}} P(X_1 = x_1, \dots, X_n = x_n) = \sum_{x_1 \text{ constant}} f_{X_1, X_2, \dots, X_n}(x_1, x_2, \dots, x_n)$$

Similarly,

Definition 2.5.2. We define marginal p.d.f as ,

$$f_{X_i}(x_i) = \int \int \int \dots \int f(x_1, x_2, \dots, x_n) dx_1 \dots dx_{i-1} dx_{i+1} \dots dx_n$$

Similarly, c.d.f follows the same rule. $F_X(x) = F(x, a, b, c, \dots)$.

Remark: Marginality and conditionality are not the same thing. They look similar, but their definitions are subtly different.

2.6 Independence

Similar to events, r.vs also can be independent,

Definition 2.6.1. Two r.vs X and Y are **independent** if, for every A and B ,

$$P(X \in A, Y \in B) = P(X \in A)P(Y \in B)$$

The definition persists for multivariate distributions.

To check Independence, we need to check the above question for every subsets A, B . Additionally, we have the theorem,

Theorem 2.6.1. Let X and Y have p.m.f $f_{X,Y}$. Then X and Y are independent only and only if ,

$$f_{X,Y}(x, y) = f_X(x)f_Y(y)$$

The definition persists for multivariate distributions.

2.7 Conditioning

Similar to events, r.v X can also have conditional distributions given that we have $Y = y$. We show the conditionality with,

Definition 2.7.1. We can show conditional distribution of X respect to Y with,

$$P(X = x|Y = y) = \frac{P(X = x, Y = y)}{P(Y = y)}$$

Moreover we can also define **conditional p.m.f** as,

Definition 2.7.2. p.m.f of X conditional respect to Y can be written as

$$f_{X|Y}(x|y) = \frac{f_{X,Y}(x, y)}{f_Y(y)}$$

2.8 Transformations of a Random Variable

In some applications, we really are interested in distributions of some function of X . We call this concept **Transformation of X**.

Definition 2.8.1. Let X be r.v with PDF/p.m.f f_X and c.d.f F_X . Let $Y = r(X)$ i.e $Y = X^2$ or $Y = \ln X$. We call $Y = r(X)$ **transformation of x** .

If Y is discrete, p.m.f is given by,

$$f_Y(y) = P(Y = y) = P(r(X) = y) = P(\{x : r(x) = y\}) = P(X \in r^{-1}(y))$$

If Y is continuous, we first calculate c.d.f and find derivative of it.

$$\begin{aligned} F_Y(y) &= P(Y \leq y) = P(r(X) < y) \\ &= P(\{x : r(x) \leq y\}) = P(A_y) \\ &= \int_{A_y} f_X(x) dx \end{aligned}$$

And the last step, $f_Y(y) = F'(y)$.

We can also generalize this concepts for Multivariate distributions, which we just increase dimensions we work with (too lazy, add this later). t

2.9 References

1. MIT

2.10 Exercises

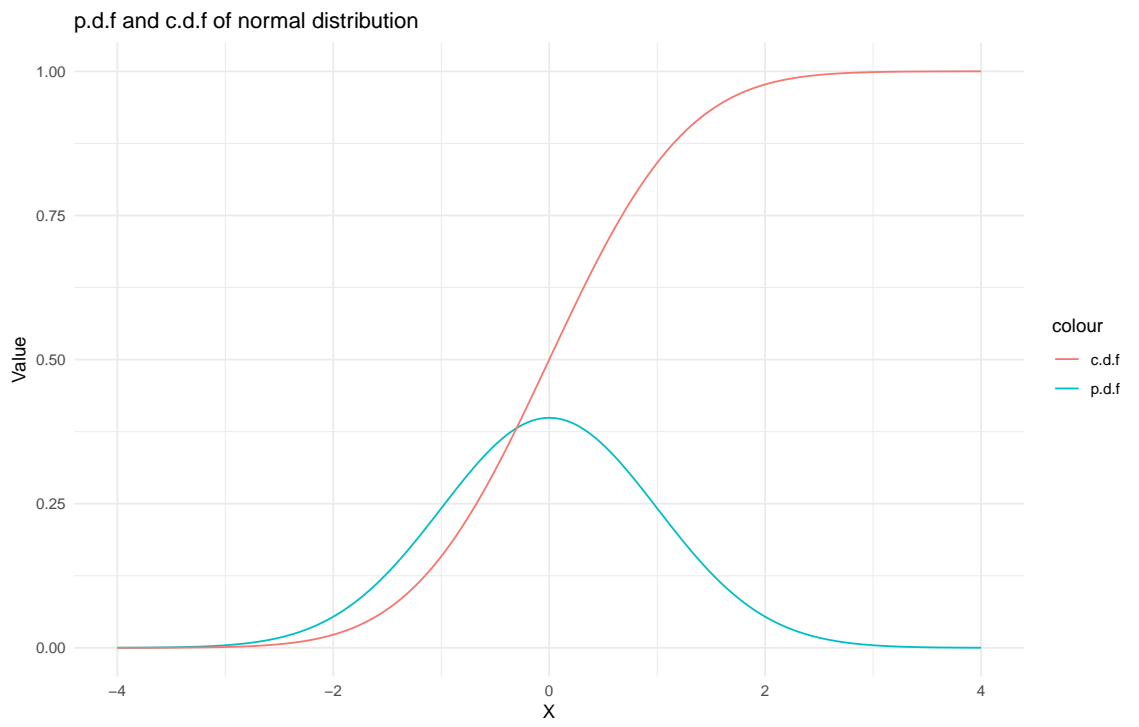
1. Suppose Z is a standard normal random variable. Use R to:

- (a) plot its p.d.f and c.d.f in the same plot

Solutions:

- (a) R already has its own unique functions that would be extremely helpful for us

```
library(tidyverse)
# R already has functions pnorm() to calculate what we need
x <- seq(-4, 4, 0.01) # to help plot c.d.f and p.d.f
y_pdf <- dnorm(x) # gives the p.d.f of Z
y_cdf <- pnorm(x) # gives the c.d.f of Z
data <- data.frame(x_value = x, y_value_pdf = y_pdf, y_value_cdf =
y_cdf)
ggplot(data) +
  geom_line(aes(x = x, y = y_value_pdf, color = "p.d.f")) +
  geom_line(aes(x = x, y = y_value_cdf, color = "c.d.f")) +
  labs(title = "p.d.f and c.d.f of normal distribution", x = "X", y
= "Value") +
  theme_minimal()
```



Chapter 3

Expectations and Invariance

3.1 Expectation of a Random Variable

The distribution of X contains all the probabilistic data we need about X . However, we need additional tools to describe these data more cleanly.

One of these tools is **Expectation**, or **Expected Value** or **Mean** of X .

Definition 3.1.1. The **expected value** of X is defined as,

$$\mathbb{E}(X) = \begin{cases} \sum xf(x) & \text{if } X \text{ is discrete} \\ \int xf(x)dx & \text{if } X \text{ is continuous} \end{cases}$$

If expected value is infinite, we say that expected value of X doesn't exist.

We can also combine both the notations into a whole generalized equation with a notation,

$$\mathbb{E}(X) = \int xdF(x) = \mu = \mu_X$$

We have discussed that $dF(x) = f(X)$. in the second chapter. **Important Note.** Expectation, by nature, is a theoretical mean of the variables we get. It is sometimes possible to get mean that you can't get in a practical settings.

By definition of probability, sum of all $f(x)$ is simply 1. Then, the above equation is weighted mean of X , which is what we wanted to convey.

Example 3.1.1. Suppose that we have a discrete r.v X to describe the probability of getting heads from tossing a coin 3 times. Let c.d.f of X be f . Then,

$$X = \begin{cases} f(0) = 1/8 \\ f(1) = 3/8 \\ f(2) = 3/8 \\ f(3) = 1/8 \end{cases}$$

Let's use our above formula to calculate $\mathbb{E}(X)$,

$$\mathbb{E}(X) = \frac{1}{8} \cdot 0 + \frac{3}{8} \cdot 1 + \frac{3}{8} \cdot 2 + \frac{1}{8} \cdot 3 = 1.5$$

This number shows that if we repeat our experiment for a very long time, the mean of the heads we got would be (or approach to) 1.5.

Observe that weighted mean is equivalent to arithmetic mean. Because getting $X = 2$ is simply getting $X = 1$ two times.

But, what if $Y = g(X)$ and we want to compute $E(Y)$? We have a theorem for that,

Theorem 3.1.1 (Law of the Unconscious Statistician). Let $Y = g(X)$. Then,

$$\mathbb{E}(Y) = \mathbb{E}(g(X)) = \int g(x)dF_X(x)$$

The general proof of this theorem is out of the scope of this book. Comparing this to original expectation equation, we can see that the only thing that changes is $g(x)$ and x , which intuitively makes sense if you think about it. In transformations, probabilities remains unchanged, while the result of probabilities gets transformed by a function.

Moreover, for a special case $g(x) = I_A(x)$, where $I_A(x) \in \{0, 1\}$ depending on $x \in A$, then,

$$\mathbb{E}(I_A(X)) = \int I_A(x)dF_X(x) = \int_A dF_X(x) = P(X \in A)$$

This means that probability is special case of expectation, which makes sense, considering probability itself is some average by definition.

Definition 3.1.2. We call n -th **raw moment** of X as

$$\mu_n = \mathbb{E}(X^n) = \int x^n dF_X(x)$$

If $\mathbb{E}(|X^k|)$ is infinite, then k^{th} moment do not exist.

We also define **k -th central moment** as moments about its mean μ i.e $\mathbb{E}[(X - \mu)^k]$. Additionally, **k -th standardized moments** as $\frac{\mathbb{E}[(X - \mu)^k]}{\sigma^n}$.

The 1st moment, the 2nd central moment, 3rd and 4rd standardized moments are called mean (expected value), **variance**, **skewness** and **kurtosis** in order. We will learn more about them in later chapters.

The moments are very useful and practical. Although there are infinitely many moments, only smaller ones are important for practical purposes. We already know the first moment and its significance.

Properties of Expectation

Theorem 3.1.2 (Non-negativity). If $X \geq 0$ is a r.v, then $\mathbb{E}(X) \geq 0$.

Proof. By definition of expectation, we have

$$\mathbb{E}(X) = \int x dF_X(x) \geq 0$$

since by definition, $dF_X(X) \geq 0$ and $x \geq 0$. □

Theorem 3.1.3 (Linearity). For **any** random variables X_1, X_2, \dots, X_n and constants a_1, a_2, \dots, a_n , we have

$$\mathbb{E}\left[\sum_i^n a_i X_i\right] = \sum_i^n a_i \mathbb{E}(X_i)$$

Proof. We will first prove the theorem for $n = 2$ with X, Y . $n = 1$ is trivial.

$$\begin{aligned}\mathbb{E}\left[a_1X_1 + a_2Y\right] &= \int (a_1x + a_2y)dF_{X,Y}(x, y) \\ &= \int (a_1x)dF_X(x) + \int (a_2y)dF_Y(y) \\ &= a_1 \int x dF_X(x) + a_2 \int y dF_Y(y) \\ &= a_1\mathbb{E}(X) + a_2\mathbb{E}(Y)\end{aligned}$$

The second line is the direct consequence of marginality. With induction, $n \geq 3$ is also true, however I will omit the solution for the sake of brevity. \square

This theorem is very useful and very practical.

Theorem 3.1.4 (multiplicity). For **independent** r.v X_1, X_2, \dots, X_n , we have

$$\mathbb{E}\left(\prod_{i=1}^n X_i\right) = \prod_{i=1}^n \mathbb{E}(X_i)$$

Proof. Similar to last one, we will use induction. $n = 1$ is trivial. For $n = 2$, let r.v be X, Y . Remember that independence has property $dF_{X,Y}(x, y) = dF_X(x) \cdot dF_Y(y)$.

$$\mathbb{E}(XY) = \int (xy)dF_{X,Y}(x, y) = \int xy dF_X(x) dF_Y(y) = \int y dF_Y(y) \int x dF_X(x) = \mathbb{E}(X)\mathbb{E}(Y)$$

For the sake of brevity, I won't show the induction part. \square

3.2 Conditional Expectation

Suppose that we want to calculate mean of X when $Y = y$. This is called conditional expectation, similar to conditional r.v and probability.

Definition 3.2.1. conditional expectation of X by $Y = y$ is given by,

$$\mathbb{E}(X|Y = y) = \int x dF_{X|Y}(x|y)$$

Note that $\mathbb{E}(X|Y)$ is a r.v itself since we don't know value of Y beforehand, or more precisely Y is a "function".

Theorem 3.2.1 (Law of total Expectations). for all r.v X and Y ,

$$\mathbb{E}[\mathbb{E}(Y|X)] = \mathbb{E}(Y) \quad \text{and} \quad \mathbb{E}[\mathbb{E}(X|Y)] = \mathbb{E}(X)$$

Proof. It is direct consequence of definition of conditional expectation and the fact that $dF(x, y) = dF(x)dF(y|x)$

$$\mathbb{E}[\mathbb{E}(Y|X)] = \int \mathbb{E}(Y|X = x)dF(x)$$

writelater \square

3.3 Variance

We have discussed about the expectation, a way of showing a property of a distribution. However, the expectation alone doesn't convey much. We have another tool called '**Variance**'. Variance, in

layman terms, describes how value of random variable varies are spread in the graph. Or in other terms, the distance between the expectation value.

We can define variance as,

Definition 3.3.1. Let X be a r.v with mean $\mu = \mathbb{E}(X)$. The Variance of X , denoted as $\mathbb{V}(X)$ or $\text{Var}(X)$ or σ^2 is the 2^{nd} central moment and is defined by,

$$\sigma^2 = \mathbb{E}[(X - \mu)^2] = \int (x - \mu)^2 dF(x)$$

We also define **standart deviation** as $\text{sd}(X) = \sqrt{\sigma^2} = \sigma$.

The standart deviation and variance convey the same information. They both represent the spread of our data. The difference between them is purely mathematical. The variance is more useful in mathematical applications, where standart deviation is very intuitive and practical. [mathisfun](#) explains it very well.

Calculating variance directly can be complicated and tedious directly sometimes. We can derive a theorem from the original definition for practical purposes.

Theorem 3.3.1. Let X be a random variable. Then,

$$\sigma^2 = E([(X - \mu)^2]) = E(X^2) - \mu^2$$

Proof. It is derived directly by algebraic manipulation and basic calculus,

$$\begin{aligned} \sigma^2 &= \int (x - \mu)^2 dF(x) \\ &= \int x^2 dF(x) - 2\mu \int x dF(x) + \mu^2 \int dF(x) \\ &= \int x^2 dF(x) - \mu^2 \\ &= \mathbb{E}(X^2) - \mu^2 \end{aligned}$$

□

3.4 Conditional Variance

Definition 3.4.1. Let $\mu = \mathbb{E}(X|Y = y)$. The **conditional variance** is defined as,

$$\mathbb{V}(X|Y = y) = \int (x - \mu)^2 dF_{X|Y}(x|y)$$

The conditional variance tells us how much of spread is left after We use $Y = y$. Reminder that $\mathbb{V}(X|Y)$ is a r.v itself since Y is a sort of "function" here.

Theorem 3.4.1 (Law of Total Variance). for any r.v X, Y , it is always true that,

$$\mathbb{V}(Y) = \mathbb{E}[\mathbb{V}(Y|X)] + \mathbb{V}(\mathbb{E}[Y|X])$$

We have stated before that $V(Y|X)$ and $E(Y|X)$ are random variables, not numbers. Therefore We compute variance and expectation of these random variables, and add them up to get the variance $V(Y)$.

3.5 Covariance and Corelation

Ley X and Y be r.v. **Covariance** and **Corelation** describes the linear relationship between X and Y .

Definition 3.5.1. If X and Y are r.v with mean μ_X , μ_Y and standart deviations σ_X , σ_Y , we define **covariance** as,

$$\text{Cov}(X, Y) = \mathbb{E} \left((X - \mu_X)(Y - \mu_Y) \right)$$

and **corelation** as,

$$\rho = \rho_{X,Y} = \frac{\text{Cov}(X, Y)}{\sigma_X \sigma_Y}$$

Notice that $\text{Cov}(X, X) = V(X)$ and $\rho_{X,X} = 1$.

Similiar to variance, calculatig covariance can be tedious. We can derive a better formula by simple algebraic manipulations,

Theorem 3.5.1. For all random variables with non-infinite means, we have

$$\text{Cov}(X, Y) = \mathbb{E}(XY) - \mathbb{E}(X)\mathbb{E}(Y)$$

Proof. Similar to Variance one, we have,

$$\begin{aligned} \text{Cov}(X, Y) &= \mathbb{E} \left((X - \mu_X)(Y - \mu_Y) \right) \\ &= \mathbb{E}(XY - X\mu_Y - Y\mu_X + \mu_X\mu_Y) \\ &= \mathbb{E}(XY) - \mu_Y\mathbb{E}(X) - \mu_X\mathbb{E}(Y) + \mu_X\mu_Y \\ &= \mathbb{E}(XY) - \mathbb{E}(X)\mathbb{E}(Y) \end{aligned}$$

□

Theorem 3.5.2. For all random variables with non-infinite means, we have,

$$-1 \leq \rho_{X,Y} \leq 1$$

Proof. It is direct consequence of Cauchy-Schwarz inequality. □

3.6 References

3.7 Exercises

Chapter 4

Statistical Inequalities

Statistical Inequalities provide a means of bounding measures and quantities, which is very useful practically and theoretically. These inequalities are used for computations on Machine learning and AI, developing new methods and techniques for practical purposes and so on. IMO, International Mathematical Olympiads, also have a “Inequalities” section with questions that heavily use some of the methods and techniques soon to be discussed here.

4.1 Probability inequalities

Probability inequalities are useful for bounding quantities that is hard to compute. Moreover, they are heavily used in the **theory of convergence**. Our first theory is,

Theorem 4.1.1. Let X be non-negative r.v and assume $\mathbb{E}(X)$ exists. Then **Markov’s inequality** states that for any $a > 0$,

$$P(X \geq a) \leq \frac{\mathbb{E}(X)}{a}$$

Proof.

$$\begin{aligned} E(X) &= \int_{-\infty}^{\infty} xf(x)dx = \int_0^{\infty} xf(x)dx \\ &= \int_0^a xf(x)dx + \int_a^{\infty} xf(x)dx \\ &\geq \int_a^{\infty} xf(x)dx \geq a \int_a^{\infty} f(x)dx \\ &= aP(X \geq a) \end{aligned}$$

□

Theorem 4.1.2. Let $\mu = E(X)$ and $\sigma^2 = V(X)$. **Chebyshev’s inequality** states that,

$$P(|X - \mu| \geq a) \leq \frac{\sigma^2}{a^2} \quad \text{and} \quad P(|Z| \geq k) \leq \frac{1}{k^2}$$

Where Z is **standard score** i.e $Z = \frac{x-\mu}{\sigma}$.

Proof. The theorem is the direct consequence of the markov inequality

$$P(|X - \mu| \geq a) = P(|X - \mu|^2 \geq a^2) \leq \frac{\mathbb{E}([X - \mu]^2)}{a^2} = \frac{\sigma^2}{a^2}$$

The second one is just a substitution. □

4.2 Expectation inequalities

Probably one of the most known inequalities of all time, that is also used in different field of mathematics, **Cauchy-Schwartz inequality** Simply states that,

Theorem 4.2.1. Cauchy-Schwarz inequality states that,

$$\mathbb{E}^2(|XY|) \leq \mathbb{E}(X^2)\mathbb{E}(Y^2)$$

Another form of this theorem is,

$$\text{Cov}^2(X, Y) \leq \sigma_X^2 \sigma_Y^2$$

This theorem is also known with its vector form. Consequently, this inequality is very useful and practical, hence there are many unique and different proofs.

This inequality is also popular on mathematical olympiads, with its familiar algebraic form.

Recall that a function f is **convex** if f is twice differentiable $f''(x) \geq 0$. Moreover f is **concave** if $-f$ is convex. Another very well known and popular inequality,

Theorem 4.2.2. Jensen's inequality states that if g is convex,

$$\mathbb{E}g(X) \geq g(\mathbb{E}X)$$

If g is concave, the inequality symbol flips. Note that we used $\mathbb{E}X$ instead of $\mathbb{E}(X)$ for asthetical purposes.

4.3 References

4.4 Exercises

Chapter 5

Convergence, CLT and LLN

5.1 Introduction

We already know the calculus definition of convergence. We say that x_n **converges** to x if, for every $\epsilon > 0$.

$$|x_n - x| < \epsilon$$

as n goes to infinity.

However, there are multiple definitions of convergence in Probability and Statistics. The core idea is simple, in layman terms, as we repeat a process for a long time, something approaches to a thing. Mathematics requires rigourosity, and hence there are multiple convergence definitions for the use case. We will learn convergence in the next section

CLT, in other words, **Central Limit theorem**, states that, under correct conditions, the distribution of **normalized version** of the sample mean **converges in distribution** to a standard normal distribution. There are multiple types of CLT.

LLN, or **Law of Large Numbers** states that under right conditions, the sample average \bar{X} **converges in probability** to the expectation $\mu = \mathbb{E}(X)$. There are also types of LLN.

5.2 Types of Convergence

There are two main, most used types of convergence in the Statistics and Probability. The weakest one is, the **convergence in distribution**

Definition 5.2.1. Let X_1, X_2, \dots be a sequence of r.v with c.d.f F_1, \dots, F_n is said to be **converging in distribution** or **converge weakly** to c.d.f F of a r.v X if

$$\lim_{n \rightarrow \infty} F_n(x) = F(x)$$

For every x that F is **continuous**. Convergence in distribution may be denoted as

$$X_n \xrightarrow{d} X$$

Concept of convergence in distribution **does not require** X_n to be close to X .

More stronger convergence, **convergence in probability** is defined as,

Definition 5.2.2. Let X_1, X_2, \dots be a sequence of r.v. X_n is said to be **converging in probability** to r.v X if for all $\epsilon > 0$,

$$\lim_{n \rightarrow \infty} P(|X_n - X| > \epsilon) = 0$$

Convergence in probability is denoted as,

$$X_n \xrightarrow{p} X$$

Visually, the convergence in probability states that, as n grows large X_n tends to be inside the ϵ brackets of X . However, **it is still possible that X_n goes out of the bound for some time**, but very unlikely.

Almost surely convergence is just a stronger version of this convergence, eliminating the ‘‘out of the bound’’ chance.

And other two less commonly used convergences,

Definition 5.2.3. Sequence $\{X_n\}$ converges **Almost surely** or **strongly** towards X if,

$$P\left(\lim_{n \rightarrow \infty} |X_n - X| = 0\right) = 1$$

Almost surely convergence is denoted as,

$$X_n \xrightarrow{a.s.} X$$

We can also rewrite the definition as, (Ferguson 1996, p. 5)

$$\lim_{n \rightarrow \infty} P\left(\sup_{m \geq n} |X_m - X| > \epsilon\right) = 0$$

Notice how this equation is similar to probability convergence.

This convergence states that, there exists a large number n that X_n **will always be** in the bounds of ϵ . There is also a stronger type of this version, called **sure convergence**. However, it is rarely used and there is no practical difference between this and the weaker version almost surely convergence. Therefore we don't talk about it.

and lastly, **convergence in mean**,

Definition 5.2.4. For a real number $r \geq 1$, X_n is said to be **converges in r -th mean** to a r.v X if,

$$\lim_{n \rightarrow \infty} \mathbb{E}(|X_n - X|^r) = 0$$

This convergence is also called L_r convergence or L^r convergence and is denoted as,

$$X_n \xrightarrow{L^r} X$$

5.3 Properties of Convergences

Here is a basic diagram showing the chain of implications, from Wikipedia.

$$\begin{array}{ccc} L^s \rightarrow & \Rightarrow_{s > r \geq 1} & L^r \rightarrow \\ & & \downarrow \\ a.s. \rightarrow & \Rightarrow & p \rightarrow \Rightarrow d \rightarrow \end{array}$$

- Almost surely convergence implies convergence in probability

$$X_n \xrightarrow{a.s.} X \Rightarrow X_n \xrightarrow{p} X$$

- Convergence in probability implies convergence in distribution

$$X_n \xrightarrow{p} X \Rightarrow X_n \xrightarrow{d} X$$

- Convergence in L^s implies convergence in L^r such that $s > r \geq 1$.

$$X_n \xrightarrow{L^s} X \Rightarrow X_n \xrightarrow{L^r} X$$

- Convergence in L^r implies convergence in probability.

$$X_n \xrightarrow{L^r} X \quad \Rightarrow \quad X_n \xrightarrow{p} X$$

All of these are intuitively, from the definition, makes sense. No proofs shall be provided.

Theorem 5.3.1. Continous mapping theorem Let $\{X_n\}$ be r.v sequence, and let X be a r.v. Let g be a continous function. then, the below is true,

$$\begin{aligned} X_n \xrightarrow{p} X &\quad \Rightarrow \quad g(X_n) \xrightarrow{p} g(X) \\ X_n \xrightarrow{d} X &\quad \Rightarrow \quad g(X_n) \xrightarrow{d} g(X) \\ X_n \xrightarrow{a.s.} X &\quad \Rightarrow \quad g(X_n) \xrightarrow{a.s.} g(X) \end{aligned}$$

Note to myself: Add addivity and multiplicative property with its proof, Slutsky's theorem.

5.4 LLN, Law of Large Numbers

There are **strong** and **weak** type of LLN, Law of large Numbers. However, both of them state the same idea, \bar{X}_n converges to μ as n goes to infinity.

Theorem 5.4.1. The weak Law of Large Numbers or shortly **WLLN** states that if X_1, \dots, X_n are i.i.d r.vs, then,

$$\bar{X}_n \xrightarrow{p} \mu$$

Proof. For the sake of simplicity, assume variance is finite. The theorem is the direct consequence of Chebysev's inequality.

$$P(|\bar{X}_n - \mu| \geq \epsilon) \leq \frac{\mathbb{V}(\bar{X}_n)}{\epsilon^2} = \frac{\mathbb{V}(X_1)}{n\epsilon^2}$$

Which right side obviously converges to 0. □

Theorem 5.4.2. The strong Law of Large Number or shortly **SLLN** states that if X_1, \dots, X_n are i.i.d r.vs and $\mu < \infty$, then,

$$\bar{X}_n \xrightarrow{a.s.} \mu$$

Proof is complex, so I will avoid giving it here.

Practically there is not much difference between WLLN and SLLN, therefore SLLN is preferred.

5.5 CLT, Central Limit Theorem

Theorem 5.5.1. Central Limit Theorem, or shortly **CLT** states that for i.i.d r.vs X_1, \dots, X_n with mean μ and variance σ^2 ,

$$\frac{\bar{X}_n - \mu}{\sqrt{\mathbb{V}(\bar{X}_n)}} \xrightarrow{d} N(0, 1)$$

5.6 References

1. <https://www.stat.cmu.edu/~larry/=stat325.01/chapter5.pdf>
2. https://en.wikipedia.org/wiki/Convergence_of_random_variables
3. <https://imai.fas.harvard.edu/teaching/files/Convergence.pdf>

4. <https://stats.stackexchange.com/questions/2230/convergence-in-probability-vs-almost-sure-convergence>
5. Lafayedemicheaux,P.,Liquet,B.(2009).UnderstandingConvergenceConcepts:AVisual-MindedandGraphical-BasedApproach.TheAmericanStatistician,63(2),173178.doi:10.1198/tas.2009.0032
6. <https://stats.stackexchange.com/questions/3734/what-intuitive-explanation-is-there-for-the-central-limit-theorem>

5.7 Exercises

Part II

Statistical Inference

Chapter 6

Statistical Models and Statistical Inference

Through the last chapter, we have talked about multiple probability distributions and their functions such as c.d.f, p.d.f, and p.m.f. However, we have assumed that we already knew the distribution and its properties. In practical world, it is not the case. We may try to find the average salary of the country, the fatality rate of a virus, and so on. Statistical inference, in shortly, is study of using the information sample we are given to deduce the characteristics of a population. Since majority of population is defined with **paramateres**, our investigation is mainly on finding, or estimating such paramteres.

6.1 Model and Inference

Assume that $X \sim N(0,1)$ and $Y \sim N(0,2)$. The X and Y don't have the same probability characteristics, more cleanly, they don't have the same distribution. However, we generalize them in a "group" called *normal distribution*. That is, there are infinitely many distribution with $N(\mu, \sigma^2)$, and to convey them in a more general way, we use **statistical models**

Definition 6.1.1. **Statistical Model** is a set of distribution such as,

$$\mathfrak{F} = \left\{ f(x; \theta) : \theta \in \Theta \right\}$$

where Θ is a **parameter space**.
There are also **nonparametric** \mathfrak{F} .

In a parametric model, if we are interested in only one parameter θ , we call it **target parameter** or **estimand**. The process of getting the estimand is called **Statistical Inference**. Majority of inferential problems are divided into three types: **estimation, confidence sets, hypothesis testing**.

6.2 Point Estimation

By convention we write the estimate of θ as $\hat{\theta}$. Since θ is constant and by definition $\hat{\theta}$ is a function, $\hat{\theta}$ is a r.v. Remark that funtion of r.vs is a r.v. In more mathematical way,

Definition 6.2.1. Let $X_1, \dots, X_n \sim F$ be i.i.d. A point estimator $\hat{\theta}$ is defined as,

$$\hat{\theta} = g(X_1, \dots, X_n)$$

We also define a very useful variable **bias** as,

$$\text{bias}(\hat{\theta}, \theta) = E(\hat{\theta}) - \theta$$

Here, θ is our target parameter, $\hat{\theta}$ is the function we use to estimate our target, or the

estimator. We usually write $\text{bias}(\hat{\theta}, \theta) = \text{bias}(\hat{\theta})$

Bias, in a literal sense, tells us the bias of the estimator we use. That is, the error that we may find when we estimate our parameter. We say that $\hat{\theta}$ is **unbiased** if,

$$\text{bias}(\hat{\theta}) = 0 \Rightarrow \mathbb{E}(\hat{\theta}) = \theta$$

We know that $\hat{\theta}$ is a r.v. We call this r.v.'s distribution as **sampling distribution**. We also define, **standard error of $\hat{\theta}$** or standard deviation,

$$\text{se} = \sqrt{\mathbb{V}(\hat{\theta})}$$

It is logical to think that the estimator should converge (with more samples) to its target value, we define such property as,

Definition 6.2.2. If a point estimator $\hat{\theta}$ converges to θ , we call that $\hat{\theta}$ is **consistent**

With bias alone, we can't characterize the quality of the estimator. Because the values of $\hat{\theta}$ may be far away than real value θ , but still be $\mathbb{E}(\hat{\theta}) = \theta$. Therefore, we also have to measure the variance in some way.

For such thing, we already have a tool,

Definition 6.2.3. The **mean square error** is defined as,

$$\text{MSE}(\hat{\theta}) = \mathbb{E}([\hat{\theta} - \theta]^2)$$

in similar fashion to the Variance definition, we can rewrite this equation as,

$$\text{MSE}(\hat{\theta}) = \text{bias}^2(\hat{\theta}) + \mathbb{V}(\hat{\theta})$$

MSE is function of both its variance and bias, hence it is a better way of showing the quality of the estimator.

Example 6.2.1. Let $X_1, \dots, X_n \sim \text{Bernoulli}(p)$. Let $\hat{p} = n^{-1} \sum_{i=1}^n X_i$. We already know that $\mathbb{E}(X) = p$ and $\mathbb{V}(X) = p(1-p)$. Our estimator is unbiased since,

$$\mathbb{E}(\hat{p}) = \frac{1}{n} \sum \mathbb{E}(X_i) = p$$

Moreover, the estimator's variance is,

$$\mathbb{V}(\hat{p}) = \mathbb{E}(\hat{p}^2) - \mathbb{E}(\hat{p})^2 = p - p^2$$

We can intuitively guess some estimators that could be effective for our purposes. But for many complex problems, it is not the case. We will learn new methods to calculate estimators in later chapters.

6.3 Properties of Point Estimation: Efficiency, Consistency, Sufficiency

Relative Efficiency

We already know that it is possible to have multiple estimators for one target parameter. We even learnt a new definition, MSE to convey the quality of such estimators. If we have two unique and unbiased estimators $\hat{\theta}_1$ and $\hat{\theta}_2$, it is logical to pick the estimator that has the lowest variance, since the lower the MSE, more efficient the estimator is. To convey such idea, we use,

Definition 6.3.1. Given two unbiased estimators $\hat{\theta}_1$ and $\hat{\theta}_2$, then **the efficiency of $\hat{\theta}_1$**

relative to $\hat{\theta}_2$, denoted as $\text{eff}(\hat{\theta}_1, \hat{\theta}_2)$, is defined as,

$$\text{eff}(\hat{\theta}_1, \hat{\theta}_2) = \frac{\mathbb{V}(\hat{\theta}_2)}{\mathbb{V}(\hat{\theta}_1)}$$

Note that if $t_1 : t_2$ is bigger than one, then it is true that $\hat{\theta}_1$ is relatively more efficient than $\hat{\theta}_2$.

Consistency

We have already talked about consistency before, we say the estimator is consistent if it converged to the target parameter,

Definition 6.3.2. The estimator $\hat{\theta}$ of θ is consistent if for any positive number ϵ ,

$$\lim_{n \rightarrow \infty} P(|\hat{\theta}_n - \theta| \leq \epsilon) = 1$$

That is, $\hat{\theta} \xrightarrow{p} \theta$

The graph (below) from latter exercises is also consistent, since visually it becomes a straight line where it equals to the target parameter.

Since consistent estimators converge to the target parameter, it is logical to think the variance also converges to 0. Think in a way that the graph of the estimator has to become straight from long wavy and curvy lines i.e it is direct consequence of convergence (real analysis stuff?). Indeed,

Theorem 6.3.1. The unbiased estimator $\hat{\theta}$ of θ is a consistent estimator if,

$$\lim_{n \rightarrow \infty} \mathbb{V}(\hat{\theta}_n) = 0$$

Sufficiency

We know that the value \bar{X} (average value) is a unbiased estimator for mean μ of X . At this point, we no longer need the sample data to estimate the μ , since we can summarize the information just with the estimator \bar{X} . But, do the \bar{X} retain all the information about X ? If it does, we call such estimator **sufficient**. That is all the sufficiency is for.

We can mathematically convey this property as conditional distribution of our sample data, given the estimator. If the distribution is dependent on our target parameter, it can't be sufficient. In more mathematical way,

Definition 6.3.3. A **statistic** is a function of data (Remark: all estimators are statistic but not all statistic are estimators). A statistic $U = t(X_1, \dots, X_n)$ of θ is sufficient if conditional distribution of X_1, \dots, X_n given U is not dependent on θ .

If conditional distribution is dependent on the target parameter, it is intuitive to think the statistic does not contain all the information.

Sufficiency is useful since it helps us to *assessing information on the entire population without the need of all the data*.

Say you get your grade on an exam and you want to know how well you did compared to your classmates. If you are given a sample mean and variance, you can do this without asking everyone's grades.

6.4 Confidence Intervals

Let's assume we are a scientist that want to measure the mean of the specific kind of mice's weight. It is unrealistic to measure **all of the mice**, hence we catch a small amount of them, probably in hundreds, measure them and gather the data in a datasheet. We **bootstrap** (we will learn what that term is in later chapters) the sample data, and find the sample mean. Now, we repeat the

bootstrapping process thousands of times, which now we have a sample mean data. Now, let's find numbers a, b such that 95% of our sample mean data resides in interval $[a, b]$. That is what **confidence interval** basically is.

Definition 6.4.1. Let X be a random vector. The $1 - \alpha$ **confidence interval** for a parameter θ is an interval $[a, b]$ and $a = a(X)$, $b = b(X)$ functions such that,

$$P(a \leq \theta \leq b) \geq 1 - \alpha$$

Note that θ is unknown constant value, while a and b are random variables.

Taking the above example, $\alpha = 0.05$, which is a mathematical standard number used majority of time. $1 - \alpha$ is called **confidence coefficient**. We also call **lower and upper confidence limits** to a and b , sometimes also denoted as θ_L and θ_U .

It is also possible to form *one sided confidence interval*, i.e.,

$$P(\theta_L \leq \theta) \geq 1 - \alpha \quad \text{or} \quad P(\theta \leq \theta_U) \geq 1 - \alpha$$

The confidence intervals may be **closed or open**. For our purpose they are indifferent.

Theorem 6.4.1 (Normal-Based Confidence Intervals). Let $z_{\alpha/2} = \Phi^{-1}(1 - \alpha/2)$. Then, there is $1 - \alpha$ chance that the below interval,

$$\hat{\theta} \pm z_{\alpha/2} \hat{se}$$

will contain the real value θ .

Proof. Let the above interval be C . Then,

$$\begin{aligned} P(\theta \in C) &= P(\hat{\theta} - z_{\alpha/2} \hat{se} < \theta < \hat{\theta} + z_{\alpha/2} \hat{se}) = \\ &= P(-z_{\alpha/2} < \frac{\hat{\theta} - \theta}{\hat{se}} < z_{\alpha/2}) \xrightarrow{d} \\ &= P(-z_{\alpha/2} < Z < z_{\alpha/2}) = 1 - \alpha \end{aligned}$$

□

6.5 Hypothesis Testing

Please read the Hypothesis chapter instead.

6.6 References

1. <https://math.stackexchange.com/questions/1767877/in-laymans-terms-what-is-the-difference-between-a-model-and-a-distribution>
2. <https://stats.stackexchange.com/questions/3911/when-are-confidence-intervals-useful>
3. <https://mason.gmu.edu/~alaemmer/bio214/sampling-distributions.pdf>

6.7 Exercises

Chapter 7

Methods of Estimation (Parametric Inference)

In later chapter we have shortly talked about the point estimation. The estimator $\hat{\theta}$ of a target parameter θ is a function of random variables of a sample, and therefore it itself is a random variable. The estimator has its own probability distribution, *sampling distribution*. We already know about *unbiased estimators* i.e $\mathbb{E}(\hat{\theta}) = \theta$ and the *consistent estimator*. In this chapter, we will learn more deeply about the mathematical properties of the point estimators. Additionally, we will learn new methods to derive estimators, since until now we listened our intuition.

7.1 Method of Moments

Until now, we have used our intuition to find estimators. For example, it is logical to think that \bar{X} would be an ideal estimator for μ of X . However, in practical world we have to generate the parametric estimators more “mathematically”. First, we introduce with a new simple definition,

Definition 7.1.1. **k-th sample moment** $\hat{\alpha}_k$ is moment of sample i.e

$$\hat{\alpha}_k = \frac{1}{n} \sum_{i=1}^n X_i^k$$

In section 3.1 we talked about **raw moments**. Raw moments convey the properties of the distribution i.e raw moments are some functions of the desired parameters. The first raw moment is the mean $\mu_1 = \mu$, the second raw moment is expression of variance $\mu_2 = \sigma^2 + \mu^2$ and so on.

The idea method of moment is we can use $\hat{\alpha}_k$ as good estimator of μ_k , and from μ_k we can derive expressions for our target parameter.

Example 7.1.1. Let $X_1, \dots, X_n \sim N(\mu, \sigma^2)$. We have that $\mathbb{E}(X) = \mu$ and $\mathbb{E}(X^2) = \sigma^2 + \mu^2$.

$$\begin{aligned}\hat{\alpha}_1 &= \hat{\mu} \\ \hat{\alpha}_2 &= \hat{\mu}^2 + \hat{\sigma}^2\end{aligned}$$

Solving the system equation will give us estimators for μ and σ .

7.2 Method of Maximum Likelihood

The method of moments are very simple and intuitive, but it is unefficient. We have a better and sophisticated method called **method of maximum likelihood**. There is a great [video](#) by Josh Starmer that explains the method very well.

Assume that we have a sample data, and we want to estimate parameters of the distribution that describes the sample data. The idea is that we find such estimator that maximize the **likelihood** of getting our sample data relative to the parameter.

Definition 7.2.1. The **Likelihood function** is defined as,

$$\mathcal{L}_n(\theta) = \prod_{i=1}^n f(X_i; \theta)$$

Also we define **log-likelihood function** as,

$$\ell_n(\theta) = \log \mathcal{L}_n(\theta)$$

At last, we define the **maximum likelihood estimator** MLE denoted by $\hat{\theta}_n$ as the value of θ that maximizes $\mathcal{L}_n(\theta)$, or better $\ell_n(\theta)$, since working with logs are easier than multiplicative functions for maximizing.

We already know that θ is an unknown constant we want to estimate. The $\mathcal{L}_n(\theta)$ describes the likelihood of each sample data, respect to θ . Since it is intuitive to maximize the likelihood (because the sample data is already happened and should be maximized), it should also estimate our value θ .

7.3 Properties of MLE

Will write later.

Chapter 8

Estimating Statistical Functions (Non-Parametric Inference)

8.1 Empirical Distribution Function, e.d.f

Definition 8.1.1. let $X_1, \dots, X_n \sim F$ be i.i.d sample with F as *c.d.f.*
We can estimate F with **empirical distribution function** or shortly **e.d.f** \widehat{F} ,

$$\widehat{F}_n(x) = n^{-1} \sum_{n=1}^n I(X \leq x)$$

Theorem 8.1.1. e.d.f \widehat{F}_n covers almost surely to F , or is **consistent** that is,

$$\widehat{F}_n(x) \xrightarrow{a.s.} F(x)$$

It also estimates F with no bias, that is,

$$\mathbb{E}(\widehat{F}_n(x)) = F(x)$$

Theorem 8.1.2. Dvoretzky–Kiefer–Wolfowitz–Massart inequality (DKW)
Let $X_1, \dots, X_n \sim F$. Then,

$$P\left(\sup_{x \in \mathbb{R}} |F(x) - \widehat{F}_n(x)| > \epsilon\right) \leq 2e^{-2n\epsilon^2} \quad \forall \epsilon > 0$$

This inequality is useful for constructing **confidence intervals** such that,

$$P\left(\widehat{F}_n(x) - \epsilon \leq F(x) \leq \widehat{F}_n(x) + \epsilon\right) = 1 - \alpha \quad \forall \epsilon = \sqrt{\frac{1}{2n} \log\left(\frac{2}{\alpha}\right)}$$

8.2 Statistical Functionals

Statistical functionals are functions of data, that is $T(F)$. Examples are mean, variance, median and so on.

Definition 8.2.1. plug-in estimator of $\theta = T(F)$ is defined as,

$$\widehat{\theta}_n = T(\widehat{F}_n)$$

That is, plug in \widehat{F}_n to estimate our statistical function $T(F)$.

Definition 8.2.2. T is **linear functional** if $T(F) = \int r(x)dF(x)$. Linear functional T satisfies linearity properties, that is,

$$T(aF + bG) = aT(F) + bT(G)$$

8.3 References

1. https://en.wikipedia.org/wiki/Empirical_distribution_function
2. https://en.wikipedia.org/wiki/Dvoretzky%E2%80%93Kiefer%E2%80%93Wolfowitz_inequality

Chapter 9

Hypothesis Testing and p-value

9.1 To-learn and write

wald test, z-test, t-test, chi, various methods and distributions.

9.2 Null and Alternate Hypothesis

The hypothesis testing is very similar to the scientific method. Scientists across the different fields use scientific method for their academical purposes. They observe, formulate a theory, experiment and test the theory. There is a similar method called **hypothesis testing** for statistical inference. First, we will introduce some notations and definitions,

Definition 9.2.1. Null Hypothesis, denoted by H_0 , is the hypothesis to be tested.
Alternate Hypothesis, denoted by H_1 , is the hypothesis contradictory to the null hypothesis. We usually try to support, since this way we could use *proof by contradiction*. If our evidence (data) favors the alternative hypothesis, we reject the null hypothesis. Formally, we wish to test,

$$H_0 : \theta \in \Theta_0 \quad \text{or} \quad H_1 : \theta \in \Theta_1$$

Where Θ_0 and Θ_1 are disjoint sets of parameter space Θ .

Example 9.2.1. Let $X_1, \dots, X_n \sim N(\mu, \sigma^2)$ with known variance σ^2 and unknown mean μ . We wish to test the hypothesis,

$$H_0 : \mu = \mu_0 \quad \text{or} \quad H_1 : \mu \neq \mu_0$$

In order to test the hypothesis, it is logical to calculate \bar{X} and compare it with μ_0 . It is reasonable to reject H_0 if \bar{X} is far away than μ_0 . But how much far away exactly?

Rejection or Critical Region is a set denoted by R to describe “how far away” the result can be. If the result is in the set R , we reject the hypothesis.

Definition 9.2.2. The **Critical Region** is defined as,

$$R = \left\{ x : T(x) > c \right\}$$

Here, c is called **critical value**, and T is a **test statistic** to help testing our hypothesis. The main question in hypothesis testing is find appropriate T and c .

In above example, \bar{X} is our T . We will learn about finding c now.

A hypothesis in the form of $\theta = \theta_0$ is called **simple** i.e θ has only one value, while in the form of $\theta > \theta_0$ or $\theta < \theta_0$ is called **composite** i.e θ has multiple values

There are 4 different possibilities we can conclude from our test. See the table for the brief introduction.

If H_0 is true, and we reject it, we call this error **type I error**. Similarly, if H_1 is true and we keep the H_0 , (or reject H_1) we call this error **type II error**. Summarize this in table,

	Retain Null	Reject Null
H_0 true	$(1 - \alpha)$	type I error (α) significance level of the test
H_1 true	type II error (β)	power of the test ($1 - \beta$)

Definition 9.2.3. The **power function** of a test with critical region R is a function defined as,

$$\beta(\theta) = P_{\theta}(X \in R)$$

that is, $\beta(\theta)$ represents the probability of rejecting H_0 if $\theta \in \Theta_0 \cup \Theta_1$ i.e H_0 or H_1 true. It is very general

The **size** of the test is defined by,

$$\alpha = \sup_{\theta \in \Theta_0} \beta(\theta)$$

that is, for θ that H_0 is true, α is the worst case scenario probability, representing **type I error** (if the test is simple).

We also define **significance level** of the test as α if its size is less than or equal to α . That is, upper bound for incorrectly rejecting H_0 . Significance level is chosen independent of data or the tests, usually as 0.05. Note in particular that both size and level don't relate to the sample.

If we are conducting pointwise test, for example $\mu = \mu_0$, then **significance level, the size of the test, and type I error** coincide, that is they are equal. However, for some special cases they may differ. Similarly, **power, 1- type II error** may also differ. So, we must be careful on our definitions when our test is not simple, that is, is composite.

9.3 p-value

The size of the test and power of the test is not that much related to our sample data. That is where **p-value** comes in.

The p-value is the probability under the null hypothesis of obtaining a real-valued test statistic at least as extreme as the one obtained, more rigorously,

Definition 9.3.1. Suppose for every $\alpha \in \{0, 1\}$ we have a rejection region R_{α} . Then **p-value** is defined as,

$$p = \inf \left\{ \alpha : T(X) \in R \right\}$$

That is, the p-value is the smallest level at which we can reject H_0 . We can calculate p with,

$$p = \sup_{\theta \in \Theta_0} P_{\theta}(T(X) \geq T(x))$$

Notice that the expression contains sample data x .

9.4 Learn and Write missing sections about Hypothesis Testing(*)

9.5 References

1. https://en.wikipedia.org/wiki/Statistical_hypothesis_test#Definition_of_terms
2. <https://stats.stackexchange.com/questions/183800/how-to-understand-the-size-of-hypothesis-tests?rq=1>
3. <https://stats.stackexchange.com/questions/299873/differences-between-p-value-level-of-significance-and-size-of-a-test>
4. <https://en.wikipedia.org/w/index.php?title=P-value&oldid=554910098>

Chapter 10

Statistical Decision Theory

Statistical Decision Theory is widely used in machine learning fields, for example training neural networks and other statistical models via computers.

There are multiple methods of generating estimators. Decision Theory helps us decide which method is more suited for our job.

The “relationship” between θ and $\hat{\theta}$ can be measured with a **loss function** $L(\theta, \hat{\theta})$. The examples are,

$$\text{Squared Error Loss} \quad (\theta - \hat{\theta})^2$$

$$\text{absolute Error Loss} \quad |\theta - \hat{\theta}|$$

$$L_p \text{Loss} \quad |\theta - \hat{\theta}|^p$$

We kind of already studied these functions before.

10.1 Risk Function

To judge the estimators, we use average of loss, or **risk**.

Definition 10.1.1. The **risk** of an estimator $\hat{\theta}$ is,

$$R(\theta, \hat{\theta}) = \mathbb{E} \left(L(\theta, \hat{\theta}) \right)$$

Notice that if our loss function is squared error,

$$R(\theta, \hat{\theta}) = \text{MSE}(\theta, \hat{\theta})$$

Part III

Statistical Models

Chapter 11

Linear Regression

Given the data with random vectors X, Y of the form

$$(\mathbf{X}, \mathbf{Y}) \sim F_{X,Y}$$

Regression is a method finding a linear function to fit properly our data. This way, we can predict the value of Y_i , or estimate in other words.

11.1 Simple Linear Regression Model

Definition 11.1.1. We can write the simplest linear regression as follows,

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$$

Here, ϵ_i is called **error variable, disturbance term, or error term**. Error term represents factors other than x that affects y . It is normal assumed that $\epsilon_i \sim N(0, \sigma^2)$. That is,

$$\mathbb{E}(\epsilon_i | X_i) = 0 \quad \wedge \quad \mathbb{V}(\epsilon_i | X_i) = \sigma^2$$

We work with conditionalities in linear regression, since the idea of the regression itself is conditional i.e Y depends on X .

The unknown variables in our model are **intercept** β_0 , **slope** β_1 and our ϵ_i . We can estimate ϵ_i with **residual** $\hat{\epsilon}_i$. In practice we would like to minimize the values of,

Definition 11.1.2. Residual Sums of Squares or shortly **r.s.s or RSS**, which measures how our model works with our data, defined as,

$$\text{RSS} = \sum_{i=1}^n \hat{\epsilon}_i^2$$

We also define **least square estimates**, that are values β_0 and β_1 such that minimizes RSS. We can derive an expression for least square estimates. In later chapters, we will also work with multivariate versions of the same topic.

Theorem 11.1.1. The least square estimates are given by,

$$\begin{aligned}\hat{\beta}_1 &= \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^n (X_i - \bar{X})^2} \\ \hat{\beta}_0 &= \bar{Y} - \hat{\beta}_1 \bar{X} \\ \hat{\sigma}^2 &= \left(\frac{1}{n-2} \right) \sum_{i=1}^n \hat{\epsilon}_i^2\end{aligned}$$

The last estimator is also called **residual standar error** or shortly RSE.

All of these estimators are unbiased. We usually denote the least square estimators as matrices i.e $\hat{\beta}^T = (\hat{\beta}_0, \hat{\beta}_1)^T$, for convenience.

Proof Sketch. Notice that

$$\text{RSS} = \sum_{i=1}^n (Y_i - \beta_0 - \beta_1 X_i)^2$$

Therefore it is enough to minimize the right expression, which is trivial via derivatives, and solving the system equations. We will later see the multivariate proof of the theorem via matrices and such. \square

11.2 Maximum Likelihood and Least Square Estimators

Assume that Y is normally distributed. It is very practical, and holds true in practice. The likelihood function is then,

$$\prod_{i=1}^n f(X_i, Y_i) = \prod_{i=1}^n f_X(X_i) \prod_{i=1}^n f_{Y|X}(Y_i|X_i) = \mathcal{L}_1 \cdot \mathcal{L}_2$$

We can use maximum likelihood method to estimate our parameters β_1, β_2 . \mathcal{L}_1 doesn't include our target parameters, so we will treat it as a constant. We wish to maximize \mathcal{L}_2 . Using the log-likelihood and our normal distribution assumption,

$$\ell(\beta_0, \beta_1, \sigma) = -n \ln(\sigma) - \frac{1}{2\sigma^2} \sum_{i=1}^n (Y_i - (\beta_0 + \beta_1 X_i))^2$$

Note that we used the p.d.f of normal distribution and the fact that $\mu = \beta_0 + \beta_1 X_i$. To maximize the log-likelihood, we have to maximize MLE of (β_0, β_1)

$$\text{RSS} = \sum_{i=1}^n (Y_i - (\beta_0 + \beta_1 X_i))^2$$

Which is the least square estimator. In fact, we have the following theorem,

Theorem 11.2.1. As long as we have the normality assumption, the least square estimator and likelihood estimator are equivalent.

11.3 Properties of Least Sqaure Estimators

We begin with the expectation and variance of least square estimators. Remember that we work with conditionals on X .

Theorem 11.3.1. The least square estimator matrix $\hat{\beta}$ is unbiased i.e

$$\mathbb{E}(\hat{\beta}|X) = \begin{pmatrix} \beta_0 \\ \beta_1 \end{pmatrix}$$

Proof Sketch. We already derived a way of writing least square estimators in the latter chap-

ter, **Theorem 12.1.1.** Using that fact and $\mathbb{E}(\epsilon|X) = 0$, it is matter of manipulating the expectation until we get our result. \square

Theorem 11.3.2. Following properties are true:

1. Consistency: $\hat{\beta} \xrightarrow{p} \beta$
2. Asymptotically normal: $\sqrt{n}(\hat{\beta} - \beta) \xrightarrow{d} N(0, 1)$
3. $1 - \alpha$ confidence intervals:

$$\beta \pm z_{\alpha/2} \widehat{\text{se}}(\beta)$$

11.4 Multivariate Regression

Suppose that, now we have multiple input variables X_1, \dots, X_k i.e

$$\mathbf{X} = \begin{pmatrix} 1 & X_{11} & \dots & X_{1k} \\ 1 & X_{21} & \dots & X_{2k} \\ \vdots & \vdots & \vdots & \vdots \\ 1 & X_{n1} & \dots & X_{nk} \end{pmatrix}$$

In this notation, we have k variables with total n datas for each variable. Similarly we define,

$$\mathbf{Y} = \begin{pmatrix} Y_1 \\ Y_2 \\ \vdots \\ Y_n \end{pmatrix} \quad \beta = \begin{pmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_k \end{pmatrix} \quad \epsilon = \begin{pmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_n \end{pmatrix}$$

Then we can write our simple regression model with matrix notations, i.e

$$\mathbf{Y} = \mathbf{X}\beta + \epsilon$$

It is very convenient to work with the matrices. The next theorem, which is the generalization of the least square estimate expressions, shows that

Theorem 11.4.1. Assume that $\mathbf{X}^T\mathbf{X}$ is invertible, then

$$\begin{aligned} \hat{\beta} &= (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{Y} \\ \mathbb{V}(\hat{\beta}|\mathbf{X}^n) &= \sigma^2(\mathbf{X}^T\mathbf{X})^{-1} \end{aligned}$$

Proof-Sketch. Let $S = \sum_{i=1}^n \hat{\epsilon}_i^2$
Assuming the inverse of $\mathbf{X}^T\mathbf{X}$ exists, we wish to minimize

$$S(\beta) = (\mathbf{Y} - \mathbf{X}\beta)^2$$

Opening the paranthesises, taking the derivative and equating to 0, (note that S is convex) proves the first part of the theorem.

Second theorem uses the definiton variance and matrix manipulation (simply put). \square

11.5 References

1. <https://stats.stackexchange.com/questions/221891/difference-between-residual-and-disturbance-epsilon>
2. <https://stats.stackexchange.com/questions/46151/how-to-derive-the-least-square-estimator-for-multiple-linear-regression>
3. <https://igpphome.ucsd.edu/~cathy/Classes/SI0223A/sio223a.chap7.pdf>

4. https://en.wikipedia.org/wiki/Proofs_involving_ordinary_least_squares#Unbiasedness_and_variance
5. <https://stats.stackexchange.com/questions/124818/logistic-regression-error-term-and-its-distribution>
6. <https://stats.stackexchange.com/questions/23479/why-do-we-assume-that-the-error-is-normally-distributed>
7. <https://stats.stackexchange.com/questions/148803/how-does-linear-regression-use-the-normal-distribution>

Chapter 12

Classification

In linear regression models, it is assumed that response variable Y is quantitative, that is, it is some real number. In practical world, Y can be qualitative e.g blood type, eye color and so on. Study of predicting qualitative responses are called **classification**.

12.1 Logistic Regression

Assume our response variable Y takes binary values i.e $Y \in \{0, 1\}$. There is a method called **logistic regression** that models the probability that Y belongs to a particular category, in this example 0 or 1, that is,

$$P(Y = 1|X) \quad \text{and} \quad P(Y = 2|X)$$

Trying linear regression model, we see that the value is not in the interval $[0, 1]$, e.g it sometimes takes negative values. To avoid this problem, we use **logistic function**.

Definition 12.1.1. Logistic function, sometimes called **sigmoid function**, is defined by formula,

$$\sigma(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{1 + e^x} = 1 - \sigma(-x)$$

Note that σ in this formula do not represent the variance, but rather the logistic function.

Using this function, we can derive a statistical model called **logistic regression**,

Definition 12.1.2. Logistic regression is defined as,

$$p(X) \equiv P(Y_i = 1|X = x) = \frac{e^{\beta_0 + \sum_{i=1}^n \beta_i x_i}}{1 + e^{\beta_0 + \sum_{i=1}^n \beta_i x_i}}$$

With simple algebraic manipulation, we can find that,

$$\ln\left(\frac{p(X)}{1 - p(X)}\right) = \beta_0 + \sum_{i=1}^n \beta_i x_i$$

The log expression in the left side is usually called **logit**, that is,

$$\text{logit}(p(X)) = \ln\left(\frac{p(X)}{1 - p(X)}\right)$$

There are subtle differences between linear regression and logistic regression. (A chapter about this?+bernoulli+error term distribution so on)

12.2 Estimating Logistic Coefficients

It is possible to generate the coefficient via MLE. Since Y is binary, the data is in binomial distribution, that is,

$$Y|X \sim \text{Bernoulli}(p)$$

therefore we wish to maximize,

$$\ell(\beta_0, \beta_1) = \prod_{i=1}^n p(x_i) \prod_{i=1}^n 1 - p(x_i)$$

Statistical softwares such as R and python libraries can easily compute the MLE. There are also multiple algorithms, one of them being **Gradient Descent**, a very useful technique used heavily in Deep Learning.

12.3 Multinomial Logistic Regression

Assume our dependent variables Y get more than 2 values. It is possible to use logistic model, with minor changes, for this purpose. This model is called **Multinomial Logistic Regression**. Let the number of classes (dependent variables) be K . The idea is choose a class as “pivot”, in this example $Y = K$, and other $K - 1$ outcomes are logistic regressed against the pivot. That is, we apply binary logistic regression multiple times, e.g (We wont write conditional part of the probabilities for the sake of the simplicity)

$$\ln \frac{P(Y_i = k)}{P(Y_i = K)} = \beta_k \cdot \mathbf{X}_i, \quad k < K$$

Note that here β_k is a matrix.

Afterwards, we use the fact that all the probability sum up to 1, then

$$P(Y_i = K) = 1 - \sum_{j=1}^{K-1} P(Y_i = j) = 1 - \sum_{j=1}^{K-1} P(Y_i = K) e^{\beta_j \mathbf{X}_i}$$

With basic algebraic manipulation, we derive that

$$P(Y_i = K) = \frac{1}{1 + \sum_{j=1}^{K-1} e^{\beta_j \mathbf{X}_i}}$$

Then other non-pivot probabilities are,

$$P(Y_i = k) = \frac{e^{\beta_k \mathbf{X}_i}}{1 + \sum_{j=1}^{K-1} e^{\beta_j \mathbf{X}_i}}$$

This statistical model works, and we estimate $K - 1$ coefficients to generate our model. However, this model is not symmetrical i.e K th class is “being ignored”. There is a better method, first we introduce **softmax function**, generalization of the logistic function

Definition 12.3.1. Softmax Function takes input matrix $\mathbf{X} = \{x_1, \dots, x_K\} \in \mathbb{R}^K$ and converts the value to probabilities,

$$\sigma(\mathbf{X})_i = \frac{e^{x_i}}{\sum_{j=1}^K e^{x_j}}$$

Intuitively, this function is also called **normalized exponential function**, because this is what basically the function does. Moreover, the sum of all the components adds up to 1.

Using the above function,

Definition 12.3.2. Softmax Multinomial Logistic Regression, for all K classes, we have

$$P(Y_i = k) = \frac{e^{\beta_k \mathbf{X}_i}}{\sum_{j=1}^K e^{\beta_j \mathbf{X}_i}}$$

12.4 Estimating Multinomial Logistic coefficients

We can either use multinomial logit, or generalized maximum likelihood for joints of coefficients, which is out of the scope of my knowledge and the book.

12.5 Linear Discriminant Analysis, l.d.a or LDA

12.6 QDA

12.7 KNN, write the above chapters after finish learning(*)

12.8 References

1. <https://stats.stackexchange.com/questions/124818/logistic-regression-error-term-and-its-distribution>
2. https://en.wikipedia.org/wiki/Multinomial_logistic_regression#

12.9 Exercises

Chapter 13

Resampling Methods

13.1 Bootstrap

Bootstrap is a technique about generating more sample data from existing data. It is used for computing confidence intervals and estimating standard errors. The idea is simple. We have sample data X_1, \dots, X_n from unknown distribution F . First we estimate F with \hat{F}_n , then we draw random resamples from \hat{F}_n multiple times and calculate our wanted statistical functions.

13.2 Basic Introduction

Definition 13.2.1. Bootstrapping

Let $X = \{X_1, \dots, X_n\}$ be our sample data from unknown distribution F . That is,

$$X_1, X_2, \dots, X_n \sim F$$

Let's assume we are interested in $T(X)$. In practical statistical problems, we need to know about the distribution of $T(X)$. We can estimate our unknown distribution F with \hat{F}_n , and draw random samples from that known distribution, that is,

$$X_1^*, X_2^*, \dots, X_n^* \sim \hat{F}_n$$

and compute T_n^* from these samples.

13.3 References

1. <https://stats.stackexchange.com/questions/26088/explaining-to-laypeople-why-bootstrapping-works>
2. <https://ocw.mit.edu/courses/14-384-time-series-analysis-fall-2013/2fdf997bca65d6ed82ba7a94f6c>

Chapter 14

Deep learning

Chapter 15

Unsupervised Learning